NewsBot 2.0 — Reflective Journal

FP_ReflectiveJournal_Solo_ITAI2373

Student: Pooja Upadhyay **Course:** ITAI 2373 — NLP **Date:** May 2025

1. Individual Contributions

I, Pooja Upadhyay, completed the NewsBot 2.0 final project independently, building all four advanced modules on top of my mid-term foundation. I was solely responsible for designing the LDA topic modeling pipeline, the extractive summarization system, the multilingual detection and translation integration, and the conversational query interface. This required understanding how each new module interacted with the existing 8-module pipeline and ensuring data flowed correctly between components.

2. Technical Integration Challenges

LDA Parameter Tuning: Finding the right number of topics (n_components) required experimentation. With 5 BBC categories I started with 5 topics, but LDA doesn't always align with human-defined categories. I found that 5 topics produced interpretable results that roughly matched the BBC taxonomy.

Summarization Quality: Extractive summarization picks existing sentences — it cannot rephrase or synthesize information. For short articles with fewer than 5 sentences, the 'summary' is nearly as long as the original. I handled this edge case by returning the full text when articles are below a length threshold.

Translation API Availability: The deep-translator library requires internet access and can hit rate limits in Google Colab. I implemented a graceful fallback that classifies the original text when translation fails, ensuring the system doesn't crash on non-English articles.

Conversational NLU Scope: Rule-based NLU has obvious limitations — it cannot understand complex or ambiguous questions. I clearly documented these limitations and implemented semantic search as a fallback for questions that don't match any rule pattern.

3. Key Findings and Business Insights

The most valuable insight from the final project was confirming that LDA topic modeling discovers the same structure as human-annotated categories — without any labels. This suggests the BBC's five categories represent genuinely distinct writing styles and vocabularies, not arbitrary editorial decisions. This has practical implications: a new news source could be categorized automatically using unsupervised methods even before any labeled training data is available.

The sentiment analysis revealed a consistent pattern: politics and business news skew negative while entertainment skews positive. This is not surprising — conflict and problems are considered more newsworthy than positive developments in serious journalism. For a media monitoring application, this means sentiment alone is not enough to detect genuine crises — baseline sentiment must be accounted for per category.

4. Professional Development Impact

This project gave me experience with the complete lifecycle of a production NLP system: from raw data cleaning through model training, evaluation, and finally user interface design. The conversational module in particular taught me how much engineering effort goes into making ML systems accessible to non-technical users — the model accuracy is less important than whether users can actually interact with the system to get value from it.

5. Comparison: NewsBot 1.0 vs 2.0

Capability	Version 1.0 (Mid-Term)	Version 2.0 (Final)
Classification	4 models compared	Linear SVM optimized + topic overlay
Themes	TF-IDF keywords only	LDA discovers hidden topics automatically
Summarization	None	Extractive 2-sentence summaries
Search	Keyword matching	Semantic cosine similarity search
Languages	English only	Detection + translation pipeline
User Interface	Code only	Natural language conversational queries

6. Future Enhancements

The most impactful future improvement would be replacing the Linear SVM classifier with a fine-tuned BERT or RoBERTa model. Transformer-based models understand context in a way bag-of-words TF-IDF cannot — for example, distinguishing 'bank robbery' (crime) from 'river bank' (geography). Based on literature, this could improve classification accuracy by 10-15 percentage points.

A Gradio or Streamlit web interface would make the system fully accessible to non-technical users without requiring Google Colab. Combined with a real-time RSS feed integration, the system could monitor live news and alert users to emerging stories matching their configured interests — the complete media intelligence product described in the business case.